

Name: Pham Minh Nghia

SID: 32507133

Problem 1:

1. Because there are 50 '5's and 50 '9's, the accuracy of an empty decision tree on the training set is 50%. If the training set is a random set then the accuracy is $\text{number_of_majority}/\text{total}$. On this particular test set, accuracy is 50% with the same above explanation.

Code:

```
#This code is part of:
#
# CMPSCI 370: Computer Vision, Spring 2021
# University of Massachusetts, Amherst
# Instructor: Subhansu Maji
#
# Mini-Project 6

import pickle
import numpy as np
import matplotlib.pyplot as plt

def scoreFeatures(x, y):
    scores = np.zeros(x.shape[:2])
    five, nine = np.zeros(x.shape[:2]), np.zeros(x.shape[:2])
    #-----
    -
    # Calculate scores (Implement this)
    #-----
    -

    # Calculate each attribute
    for i in range(len(x[:, :, 1])):          # loop through each pixel feature
        for j in range(len(x[i, :, 1])):
            for img in range(len(x[i, j])):    # loop through each image
                if x[i, j, img] == 0:          # 1 is on and 0 is off
                    if y[img] == 5:            # 5 is 'hated'
                        five[i, j] += 1
                else:
                    if y[img] == 9:            # 9 is 'liked'
                        nine[i, j] += 1

    scores = five + nine
    plt.imshow(scores, cmap='gray')
    plt.axis('off')
    plt.show()
    return scores
```

```

def highest_score(scores):
    r_i, r_j = 0, 0
    highest = 0
    for i in range(scores.shape[0]):
        for j in range(scores.shape[1]):
            check = highest
            highest = max(highest, scores[i, j])
            if highest != check:
                r_i, r_j = i, j
    print(r_i, r_j, highest)
    return r_i, r_j

def decisiontree_test(x, y, i, j):
    accuracy = 0
    predict = []
    for img in range(len(x[i, j])): # loop through each image
        if x[i, j, img] == 0:
            predict.append(5)
        else:
            predict.append(9)

    for k in range(len(predict)):
        if predict[k] == y[k]:
            accuracy += 1
    print("Accuracy =", accuracy/len(x[i, j]) * 100, '%')
    return None

def depth2(x, y, x_test, y_test):
    scores_p = scoreFeatures(x, y)
    i_p, j_p = highest_score(scores_p)
    sub_l, sub_r = [], [] # Assign left dataset and right dataset
    label_l, label_r = [], [] # left and right label

    for img in range(len(x[i_p, j_p])): # loop through each image
        if x[i_p, j_p, img] == 0: # pixel=0 go to left
            sub_l.append(x[:, :, img])
            label_l.append(y[img])
        else: # pixel=1 go to right
            sub_r.append(x[:, :, img])
            label_r.append(y[img])

    # convert to numpy array
    sub_l = np.array(sub_l)
    sub_r = np.array(sub_r)

    # calculate scores for left and right datasets
    scores_l = scoreFeatures(sub_l.transpose(1, 2, 0), label_l)
    scores_r = scoreFeatures(sub_r.transpose(1, 2, 0), label_r)
    i_l, j_l = highest_score(scores_l)
    i_r, j_r = highest_score(scores_r)

    # Prediction
    predict = []
    for img in range(len(x_test[i_p, j_p])): # loop through each image
        if x_test[i_p, j_p, img] == 0:

```

```

        if x_test[i_l, j_l, img] == 0:
            predict.append(5)
        else:
            predict.append(9)
    else:
        if x_test[i_r, j_r, img] == 0:
            predict.append(5)
        else:
            predict.append(9)

# Calculate accuracy
accuracy = 0
for i in range(len(predict)):
    if predict[i] == y_test[i]:
        accuracy += 1
print("Accuracy =", accuracy / len(x_test[i_p, j_p]) * 100, '%')

def main():
    data = pickle.load(open('data_binarized.pkl', 'rb'))
    scores = scoreFeatures(data['train']['x'], data['train']['y'])
    i, j = highest_score(scores)
    decisiontree_test(data['test']['x'], data['test']['y'], i, j)

    depth2(data['train']['x'], data['train']['y'], data['test']['x'],
data['test']['y'])
    #-----
    -
    # Your implementation to answer questions on Decision Trees
    #-----
    -

    return

if __name__ == "__main__":
    main()


```

Result:

Depth = 1

Highest Value: [11 18]= 163.0

Accuracy = 82.5 %





Depth = 2

Parent: [11 18]= 163.0

Left: [12 20]= 100.0

Right: [13 17]= 74.0

Accuracy = 90.5 %



Problem 2:

Code:

```
#This code is part of:
#
#  CMPSCI 370: Computer Vision, Spring 2021
#  University of Massachusetts, Amherst
#  Instructor: Subhransu Maji
#
#  Mini-Project 6

import pickle
import numpy as np
import matplotlib.pyplot as plt

def visualizeNeighbors(imgs, topk_idx, topk_distances, title):
    """
    Visualize the query image as well as its nearest neighbors
    Input:
        imgs: a list or numpy array, with length k+1.
              imgs[0] is the query image with shape h x w
              imgs[k] is k-th nearest neighbor image
        topk_idx: a list or numpy array, with length k+1.
                 topk_idx[k] is the index in training set of the k-th nearest image
                 topk_idx[0] is the query image index in the test set
        topk_distances: a list or numpy array, with length k+1.
                       topk_distances[k] is the distance of the k-th nearest image to the query
                       topk_distances[0] is 0
    """
    n = len(imgs)
    fig, axs = plt.subplots(1, n, figsize=(2 * n, 3))
    fig.suptitle(title)
    for k in range(n):
        if k == 0:
            ax_title = 'query: test_idx=%d' % topk_idx[0]
        else:
            ax_title = '%d: idx=%d, d=%.2e' % (k, topk_idx[k], topk_distances[k])
        axs[k].set_title(ax_title)
        axs[k].imshow(imgs[k], cmap='gray')
        axs[k].axis('off')
    fig.tight_layout()
    plt.show()

    return

def knn(visualize):
    data = pickle.load(open('data_binarized.pkl', 'rb'))

    # distances = np.zeros((len(data['test']['y']), len(data['train']['y'])))

    # List of dictionary
    # element i-th is a dictionary with key [i-th, train img]
    # dict keys: [test img, train img]
    # dict value: Euclidean distance
    distances = [{}] * 200
    #-----
```

```

-
# Your implementation to calculate and sort distances
#-----

x_train = data['train']['x']
x_test = data['test']['x']
for i in range(len(data['test']['y'])):      # loop through each image
    diction = {}
    for j in range(len(data['train']['y'])):
        # Apply Euclidean formula
        diction[i, j] = np.sqrt(np.sum(np.square(x_test[:, :, i] -
x_train[:, :, j])))
    # Sort dictionary by values
    diction = dict(sorted(diction.items(), key=lambda item: item[1]))
    distances[i] = diction

if visualize:
    k = 5
    imgs = np.random.randint(2, size=(k+1, 28, 28))
    topk_idxes = [0] * (k+1)
    topk_distances = [0] * (k+1)
    for test_i in [10, 20, 110, 120]:
        topk_idxes[0] = test_i
        # Assign query img as test_i
        imgs[0] = x_test[:, :, test_i]
        #-----
        # Prepare imgs, topk_idxes and topk_distances
        #-----
        i = 0
        for key in distances[test_i]:
            i += 1      # Only go through k nearest neighbor
            imgs[i] = x_test[:, :, key[1]]      # key[1] is index for training
img of k-th neighbor
            topk_idxes[i] = key[1]
            topk_distances[i] = distances[test_i][key]
            if i >= k:
                break

        visualizeNeighbors(imgs, topk_idxes, topk_distances,
            title='Test img %d: Top %d Neighbors' % (test_i, k))

k_list = [1, 3, 5, 7, 9]
accuracy_list = [0.0] * len(k_list)
#-----

-
# Your implementation to calculate knn accuracy
#-----

-
for k, acc in zip(k_list, accuracy_list):
    # Go through all test images
    for i in range(len(data['test']['y'])):
        five, nine, j = 0, 0, 0
        for key in distances[i]:
            j += 1      # Only check k nearest neighbor

            # Finding majority
            if data['train']['y'][key[1]] == 5:

```

```

        five += 1
    else:
        nine += 1
    if j >= k:
        break

    if five >= nine:        # 5 is majority
        if data['test']['y'][i] == 5:
            acc += 1
    else:                   # 9 is majority
        if data['test']['y'][i] == 9:
            acc += 1
    # divide by total image (200) to get probability
    acc /= len(data['test']['y'])
    print('k=%d: accuracy=%.2f%%' % (k, acc * 100))

return

if __name__ == "__main__":
    knn(visualize=True)

```

Result:

k=1: accuracy=98.00%

k=3: accuracy=97.50%

k=5: accuracy=97.50%

k=7: accuracy=97.50%

k=9: accuracy=98.00%


Test img 10: Top 5 Neighbors

query: test_idx=10 1: idx=75,d=9.85e+00 2: idx=82,d=9.85e+00 3: idx=7,d=1.01e+01 4: idx=87,d=1.04e+01 5: idx=94,d=1.06e+01




Test img 20: Top 5 Neighbors

query: test_idx=20 1: idx=30,d=7.35e+00 2: idx=17,d=7.55e+00 3: idx=85,d=8.19e+00 4: idx=84,d=8.25e+00 5: idx=48,d=8.43e+00




Test img 110: Top 5 Neighbors

query: test_idx=110 1: idx=186,d=7.75e+002: idx=147,d=7.81e+003: idx=175,d=7.94e+004: idx=167,d=8.00e+005: idx=190,d=8.00e+00



Test img 120: Top 5 Neighbors

query: test_idx=120 1: idx=143,d=5.83e+002: idx=139,d=6.00e+003: idx=164,d=6.40e+004: idx=104,d=6.78e+005: idx=138,d=6.86e+00



Problem 3:

Code:

```
#This code is part of:
#
#  CMPSCI 370: Computer Vision, Spring 2021
#  University of Massachusetts, Amherst
#  Instructor: Subhransu Maji
#
#  Mini-Project 6

import pickle
import numpy as np
import matplotlib.pyplot as plt

def softmax(z):
    return 1.0/(1+np.exp(-z))

def linearTrain(x, y):
    #Training parameters
    maxiter = 50
    lamb = 0.01
    eta = 0.01

    #Add a bias term to the features
    x = np.concatenate((x, np.ones((1, x.shape[1]))), axis=0)

    class_labels = np.unique(y)
    num_class = class_labels.shape[0]
    assert(num_class == 2) # Binary labels
    num_feats = x.shape[0]
    num_data = x.shape[1]
```



```

true_prob = np.zeros(num_data)
true_prob[y == class_labels[0]] = 1

#Initialize weights randomly
model = {}
model['weights'] = np.random.randn(num_feats)*0.01
# print('w', model['weights'].shape)
#Batch gradient descent
verbose_output = False
for it in range(maxiter):
    prob = softmax(model['weights'].dot(x))
    delta = true_prob - prob
    gradL = delta.dot(x.T)
    model['weights'] = (1 - eta*lamb)*model['weights'] + eta*gradL
model['classLabels'] = class_labels

return model

def linearPredict(model, x):
    #Add a bias term to the features
    x = np.concatenate((x, np.ones((1, x.shape[1]))), axis=0)

    prob = softmax(model['weights'].dot(x))
    ypred = np.ones(x.shape[1]) * model['classLabels'][1]
    ypred[prob > 0.5] = model['classLabels'][0]

    return ypred

def testLinear():
    #-----
    --
    # Your implementation to answer questions on Linear Classifier
    #-----
    --

    data = pickle.load(open('data_binarized.pkl', 'rb'))
    x, y = data['train']['x'], data['train']['y']

    # Reshape from HxWx200 to Nx200, N = H*W
    x = np.reshape(x, (len(data['train']['x'][:, :, 1]) *
len(data['train']['x'][:, :, 1]),
len(data['train']['y'])))
    model = linearTrain(x, y)

    # Reshape testing data
    x_test = data['test']['x']
    x_test = np.reshape(x_test, (len(data['test']['x'][:, :, 1]) *
len(data['test']['x'][:, :, 1]),
len(data['test']['y'])))
    # Predict test data using trained model
    ypred = linearPredict(model, x_test)
    acc = 0
    for i in range(len(data['test']['y'])):
        # Compare predict to label of testing
        if ypred[i] == data['test']['y'][i]:

```

```

        acc += 1
    acc /= len(data['test']['y'])
    print(acc*100, '%')

    # Visualization part
    dimension = model["weights"][:len(model["weights"])-1]
    w = np.reshape(dimension, (int(np.sqrt(len(dimension))),
int(np.sqrt(len(dimension)))))
    wp = np.clip(w, 0, None)
    wn = np.clip(w, None, 0)
    plt.subplot(131)
    plt.title('Positive Weights')
    plt.imshow(wp, cmap='gray')
    plt.subplot(133)
    plt.title('Negative Weights')
    plt.imshow(wn, cmap='gray')
    plt.show()

    return None

if __name__ == "__main__":
    testLinear()

```

Result:

Accuracy= around 98%

